

Problem 1

- (a) Assume to the contrary we have some \mathbf{w}, b, ξ that minimizes the objective function under the given constraints s.t $\exists \xi_k (\xi_k < 0)$. Define a new $\xi'_n := \max(0, \xi_n)$. Because $\xi_k^2 > \xi'_k{}^2$, it follows

$$\frac{1}{2} \|\mathbf{w}\|_2^2 + C \xi_k^2 + C \sum_{\substack{n=1 \\ n \neq k}}^N \xi_n^2 > \frac{1}{2} \|\mathbf{w}\|_2^2 + C \sum_{n=1}^N \xi'_n{}^2.$$

It suffices to show ξ' satisfies $t_n(\mathbf{w}^\top \phi(x_n) + b) \geq 1 - \xi'_n$.

$t_n(\mathbf{w}^\top \phi(x_n) + b) \geq 1 - \xi_n \geq 1 - \xi'_n$ because $\xi'_n \geq \xi_n (\forall n)$, so ξ' also satisfies the given constraints. Hence, we obtain a contradiction because any solution containing $\xi_k < 0$ can't be a minimizer, so the $\xi_n \geq 0 \quad \forall n$ doesn't affect the optimal solution.

- (b) C is a regularization parameter that controls the tradeoff between maximizing the margin and maximizing the classification error. Increasing C leads to a smaller margin, but fewer misclassifications on the training set. This can lead to greater variance and overfitting.

Problem 2

Assume we have some \mathbf{w}_0 and b_0 s.t all points satisfy the constraint in Eq (7.5) and maximizes Eq (7.3). It follows the hyperplane $\mathbf{w}_0^\top x + b_0 = 0$ is the optimal classification margin. Suppose the constraint in Eq (7.5) becomes

$$t_n(\mathbf{w}^\top \phi(x_n) + b) \geq \gamma \quad \forall n$$

If we perform the change of variables $\mathbf{w}_0 \rightarrow \gamma \mathbf{w}_0$ and $b_0 \rightarrow \gamma b_0$, the new constraint will be satisfied and Eq (7.3) will still be maximized. The new hyperplane $\gamma \mathbf{w}_0^\top x + \gamma b_0 = 0$ is the optimal classification margin. However, the new hyperplane is equivalent to the old hyperplane ($\gamma \mathbf{w}_0^\top x + \gamma b_0 = 0$ iff $\mathbf{w}_0^\top x + b_0 = 0$). Moreover, the minimum distance to the hyperplane remains the same ($\frac{t_n(\gamma \mathbf{w}^\top \phi(x_n) + \gamma b)}{\|\gamma \mathbf{w}\|} = \frac{\gamma t_n(\mathbf{w}^\top \phi(x_n) + b)}{\gamma \|\mathbf{w}\|} = \frac{t_n(\mathbf{w}^\top \phi(x_n) + b)}{\|\mathbf{w}\|}$).

Problem 3

Preprocessing: The steps I took to preprocess the MNIST dataset were as follows: applied scikit-learn's Standard Scaler to the entire dataset, split into training (80%) and test (20%), fit a PCA to the training set that preserved 95% of the explained variance, applied PCA transformation to the train and test set. K-Nearest Neighbors tested number of nearest neighbors [5, 10, 15, 20]. Used scikit-learn's GridSearchCV with 3 folds to train model. 5 neighbors performed the best.

Across the board model performed very well on the test set. Performed best overall. Class-8 only had a 91% recall most often misclassifying 8s as 5s (falsely

classified as 5s). Otherwise, all metrics in the mid to high 90%.

Accuracy: 95%, Precision: 95%, Recall: 95%, f-1 score: 95%

Logistic Regression tested the regularization value C (inverse of regularization strength) [0.001, 0.01, 0.1, 1, 10, 100]. Used scikit-learn's GridSearchCV with 3 folds to train model. $C = 1$ performed the best.

Model performed pretty well on the test set. Model did not perform as well classifying 3s, 5s, 8s, and 9s. 3, 5, and 8 had slightly lower recalls than the other classes, and 8 and 9 had slightly lower precisions.

Accuracy: 92%, Precision: 92%, Recall: 92%, f-1 score: 92%

Decision tree tested max depth of the tree [*None*, 5, 10, 15, 20]. Used scikit-learn's GridSearchCV with 5 folds to train model. max depth=15 performed the best.

Model performed the worst relative to other models. Did not perform well classifying 3s, 5s, 8s, and 9s as with Logistic Regression. 3, 5, 8, and 9 had lower recalls whereas 5, 8, and 9 had lower precisions.

Accuracy: 83%, Precision: 83%, Recall: 83%, f-1 score: 83%

Linear svm classification tested the regularization value C [0.001, 0.01, 0.1, 1]. Used scikit-learn's GridSearchCV with 2 folds to train model. $C = 0.01$ performed the best.

Note: $C = 0.001$ took about 10% the time it took to fit as the other tested values, and the accuracies were all within less than 0.5%. Like previous models, performs slightly worse on 3s, 5s, 8s, and 9s (both accuracy and precision).

Accuracy: 91%, Precision: 91%, Recall: 91%, f-1 score: 91%

Svm classification tested the different kernels linear, poly, rbf, and sigmoid. Used scikit-learn's GridSearchCV with 5 folds to train model and $C = 0.001$. The linear kernel performed the best.

Other kernels don't perform nearly as well. Linear kernel performs pretty well across the board. Slightly worse recall on 5s.

Used $C = 0.001$ instead of $C = 0.01$ to reduce training time. Tested the hyperparameters C , kernel individually to reduce training time as well.

Accuracy: 94%, Precision: 94%, Recall: 94%, f-1 score: 94%

K-Nearest Neighbors performed the best in comparison to other models and didn't take very long to fit/score.

Classification on MNIST Dataset

In the question, you are asked to write code for classifying image data using various classifiers. The [MNIST dataset](#) is a database of 70,000 28×28 pixel grayscale images of handwritten single digits.

A Few Task Reminders:

- Import required libraries.
- Load the data and split into train, validation, and test sets. You can also perform k-fold cross-validation on the train set for better performance estimates and nested cross-validation for hyperparameter tuning.
- Perform any required data pre-processing.
- Train K-NN, Logistic Regression, Decision Trees, and SVM on the data.
- Make predictions, evaluate, and compare the models. Generate confusion matrices and classification reports.
- Summarize your findings and make sure you sufficiently document your code.

Explore the different classifiers listed above and perform hyperparameter tuning as follows:

- For K-NN, explore the effect of varying the number of nearest neighbors
- For Logistic Regression, explore the effect of varying regularization parameter
- For Decision Trees, explore the effect of varying the max depth of the tree
- For SVM, explore the effect of varying the penalty parameter and kernel function

Note. It is intentional that this problem assignment extends outside of what we have covered in class (i.e. text data pre-processing) and encourages more independent learning and exploration with ML hands-on experience and applications. I hope you would have fun with these type of questions and that they are not very stressful. Also, feedback is welcomed and encouraged!

```
In [2]: import numpy as np
import matplotlib.pyplot as plt
from sklearn.datasets import fetch_openml
```

```
In [12]: # Fetch MNIST data (might take some time)
mnist = fetch_openml('mnist_784')

X = mnist.data.astype('float32')
y = mnist.target.astype('int64')

# Normalize the data
X /= 255.0
```

```
In [144... from sklearn.neighbors import KNeighborsClassifier
from sklearn.linear_model import LogisticRegression
from sklearn.tree import DecisionTreeClassifier
from sklearn.svm import SVC, LinearSVC
from sklearn.model_selection import GridSearchCV, train_test_split
```

```

from sklearn.preprocessing import StandardScaler
from sklearn.decomposition import PCA
from sklearn.metrics import classification_report, confusion_matrix
import pickle

```

```

In [76]: scaler=StandardScaler()
X_scaled=scaler.fit_transform(X)

```

```

In [77]: X_train,X_test,y_train,y_test=train_test_split(X_scaled,y,test_size=0.2)

```

```

In [78]: pca=PCA(0.95)
pca.fit(X_train)

```

```

Out[78]: PCA(n_components=0.95)

```

```

In [79]: X_train_pca=pca.transform(X_train)
X_test_pca=pca.transform(X_test)

```

```

In [164... search_space_k={
    'n_neighbors' : [5,10,15,20]
}
search_space_l={
    'C' : [0.001, 0.01, 0.1, 1, 10, 100, 1000]
}
search_space_t={
    'max_depth': [None, 5,10,15,20],
}
search_space_s={
    'kernel' : ['linear', 'poly', 'rbf', 'sigmoid']
}
search_space_lin_svc={
    'C' : [0.001, 0.01, 0.1, 1]
}

```

```

In [81]: knn=KNeighborsClassifier()
clf_knn=GridSearchCV(knn,search_space_k,cv=3,n_jobs=-1)
clf_knn.fit(X_train_pca,y_train)

```

```

Out[81]: GridSearchCV(cv=3, estimator=KNeighborsClassifier(), n_jobs=-1,
    param_grid={'n_neighbors': [5, 10, 15, 20]})

```

```

In [98]: clf_knn.best_params_

```

```

Out[98]: {'n_neighbors': 5}

```

```

In [84]: clf_knn.cv_results_

```

```

Out[84]: {'mean_fit_time': array([0.92393788, 0.92772937, 0.91927274, 1.54400261]),
    'std_fit_time': array([0.46382901, 0.46329492, 0.45987636, 0.20259877]),
    'mean_score_time': array([161.5054493 , 170.73705546, 161.42259161, 90.773199

```

```

]),
'std_score_time': array([18.22641628,  4.96037859, 18.59974226, 53.06128924]),
'param_n_neighbors': masked_array(data=[5, 10, 15, 20],
                                   mask=[False, False, False, False],
                                   fill_value='?',
                                   dtype=object),
'params': [{'n_neighbors': 5},
            {'n_neighbors': 10},
            {'n_neighbors': 15},
            {'n_neighbors': 20}],
'split0_test_score': array([0.94166176, 0.93817968, 0.93437617, 0.9305191 ]),
'split1_test_score': array([0.94476884, 0.94048321, 0.93651899, 0.93314405]),
'split2_test_score': array([0.93887282, 0.93405122, 0.92997964, 0.92714026]),
'mean_test_score': array([0.94176781, 0.93757137, 0.93362493, 0.9302678 ]),
'std_test_score': array([0.00240821, 0.00266085, 0.00272201, 0.00245747]),
'rank_test_score': array([1, 2, 3, 4], dtype=int32)}

```

```
In [87]: predicted_knn=clf_knn.predict(X_test_pca)
```

```
In [121... print(classification_report(y_test,predicted_knn))
```

	precision	recall	f1-score	support
0	0.97	0.98	0.98	1362
1	0.97	0.99	0.98	1624
2	0.96	0.95	0.95	1390
3	0.93	0.96	0.95	1415
4	0.96	0.95	0.95	1361
5	0.93	0.92	0.93	1271
6	0.96	0.97	0.96	1391
7	0.94	0.96	0.95	1413
8	0.97	0.91	0.94	1372
9	0.93	0.94	0.94	1401
accuracy			0.95	14000
macro avg	0.95	0.95	0.95	14000
weighted avg	0.95	0.95	0.95	14000

```
In [125... confusion_matrix(y_test,predicted_knn)
```

```

Out[125... array([[1339,  0,  2,  2,  1,  3, 14,  0,  0,  1],
 [  0, 1607,  2,  1,  0,  2,  6,  4,  1,  1],
 [ 13,  6, 1316, 18,  4,  0,  8, 17,  5,  3],
 [  3,  4,  8, 1358,  1, 15,  1, 12,  5,  8],
 [  0,  9, 13,  2, 1289,  3,  5,  9,  2, 29],
 [  5,  2,  5, 39,  5, 1174, 20,  2,  9, 10],
 [  8,  5,  6,  1,  7,  10, 1349,  0,  5,  0],
 [  0,  8,  8,  4, 10,  0,  0, 1353,  3, 27],
 [  8, 18, 11, 21,  5, 44,  2,  7, 1243, 13],
 [  4,  3,  6, 13, 22,  6,  0, 33,  2, 1312]])

```

```

In [107... with open('model_knn.pkl','wb') as f:
            pickle.dump(clf_knn,f)

#with open('model_knn.pkl','rb') as f:
#    clf2=pickle.load(f)

```

```
In [173... logr=LogisticRegression(solver='sag',max_iter=1000,tol=1e-3)
clf_logr=GridSearchCV(logr,search_space_l,cv=3,n_jobs=-1)
clf_logr.fit(X_train_pca,y_train)
```

```
Out[173... GridSearchCV(cv=3,
              estimator=LogisticRegression(max_iter=1000, solver='sag',
                                             tol=0.001),
              n_jobs=-1, param_grid={'C': [0.001, 0.01, 0.1, 1, 10, 100, 1000]})
```

```
In [174... clf_logr.best_params_
```

```
Out[174... {'C': 1}
```

```
In [175... clf_logr.cv_results_
```

```
Out[175... {'mean_fit_time': array([ 25.23615702,  80.89298693, 120.65007218, 127.03134791,
                             121.07482306, 113.22952151, 106.64190173]),
            'std_fit_time': array([0.76727742, 2.92715113, 7.1853959 , 7.42636872, 3.819667
43,
                                4.16910999, 5.81419394]),
            'mean_score_time': array([0.03298338, 0.04689233, 0.03929575, 0.03024785, 0.016
72188,
                                0.01860929, 0.0159061 ]),
            'std_score_time': array([0.01025135, 0.01655264, 0.00589005, 0.01312707, 0.0012
6374,
                                0.00218765, 0.00259891]),
            'param_C': masked_array(data=[0.001, 0.01, 0.1, 1, 10, 100, 1000],
                                     mask=[False, False, False, False, False, False, False],
                                     fill_value='?',
                                     dtype=object),
            'params': [{'C': 0.001},
                       {'C': 0.01},
                       {'C': 0.1},
                       {'C': 1},
                       {'C': 10},
                       {'C': 100},
                       {'C': 1000}],
            'split0_test_score': array([0.90989447, 0.91825146, 0.92098355, 0.92087641, 0.9
2098355,
                                0.92092998, 0.92092998]),
            'split1_test_score': array([0.91123373, 0.91964429, 0.92039428, 0.92066213, 0.9
2060856,
                                0.92055499, 0.92055499]),
            'split2_test_score': array([0.90849673, 0.91808636, 0.91942569, 0.91969356, 0.9
1958641,
                                0.91963999, 0.91963999]),
            'mean_test_score': array([0.90987498, 0.9186607 , 0.92026784, 0.9204107 , 0.920
39284,
                                0.92037499, 0.92037499]),
            'std_test_score': array([0.00111746, 0.00069876, 0.00064225, 0.00051459, 0.0005
9042,
                                0.0005418 , 0.0005418 ]),
            'rank_test_score': array([7, 6, 5, 1, 2, 3, 3], dtype=int32)}
```

```
In [176... predicted_logr=clf_logr.predict(X_test_pca)
```

In [177...

```
print(classification_report(y_test,predicted_logr))
```

	precision	recall	f1-score	support
0	0.96	0.97	0.96	1362
1	0.95	0.97	0.96	1624
2	0.93	0.90	0.91	1390
3	0.91	0.89	0.90	1415
4	0.92	0.95	0.94	1361
5	0.90	0.87	0.88	1271
6	0.94	0.95	0.94	1391
7	0.93	0.94	0.93	1413
8	0.89	0.89	0.89	1372
9	0.90	0.91	0.90	1401
accuracy			0.92	14000
macro avg	0.92	0.92	0.92	14000
weighted avg	0.92	0.92	0.92	14000

In [178...

```
confusion_matrix(y_test,predicted_logr)
```

```
Out[178...] array([[1318, 1, 2, 4, 2, 8, 12, 6, 6, 3],
      [ 0, 1574, 6, 8, 2, 6, 3, 8, 15, 2],
      [ 9, 15, 1250, 16, 14, 8, 16, 19, 36, 7],
      [ 7, 7, 38, 1258, 3, 39, 8, 10, 28, 17],
      [ 0, 5, 6, 2, 1290, 3, 12, 3, 5, 35],
      [18, 6, 6, 43, 14, 1102, 30, 3, 40, 9],
      [15, 4, 10, 1, 14, 18, 1318, 5, 5, 1],
      [ 0, 8, 10, 4, 11, 2, 0, 1330, 3, 45],
      [10, 36, 10, 28, 7, 34, 6, 5, 1215, 21],
      [ 3, 7, 5, 20, 39, 7, 0, 43, 7, 1270]])
```

In [179...

```
with open('model_logr.pkl','wb') as f:
    pickle.dump(clf_logr,f)
```

In [110...

```
decision_tree=DecisionTreeClassifier()
clf_dct=GridSearchCV(decision_tree,search_space_t,n_jobs=-1)
clf_dct.fit(X_train_pca,y_train)
```

Out[110...

```
GridSearchCV(estimator=DecisionTreeClassifier(), n_jobs=-1,
              param_grid={'max_depth': [None, 5, 10, 15, 20]})
```

In [111...

```
clf_dct.best_params_
```

Out[111...

```
{'max_depth': 15}
```

In [112...

```
clf_dct.cv_results_
```

Out[112...

```
{'mean_fit_time': array([37.50029001, 10.53001604, 18.43235664, 24.73020592, 26.
58777823]),
 'std_fit_time': array([1.01606573, 0.06931991, 0.10673776, 0.07807035, 0.250087
74]),
 'mean_score_time': array([0.01499977, 0.01235256, 0.00611401, 0.00481806, 0.005
67966]),
 'std_score_time': array([0.0073633 , 0.0039344 , 0.00179303, 0.0007732 , 0.0005
```

```

2917])),
'param_max_depth': masked_array(data=[None, 5, 10, 15, 20],
                                mask=[False, False, False, False, False],
                                fill_value='?',
                                dtype=object),
'params': [{'max_depth': None},
            {'max_depth': 5},
            {'max_depth': 10},
            {'max_depth': 15},
            {'max_depth': 20}],
'split0_test_score': array([0.81169643, 0.58803571, 0.79258929, 0.820625 , 0.8
1482143]),
'split1_test_score': array([0.81214286, 0.59535714, 0.80125 , 0.82294643, 0.8
1410714]),
'split2_test_score': array([0.819375 , 0.59642857, 0.80723214, 0.82705357, 0.8
2133929]),
'split3_test_score': array([0.81803571, 0.59678571, 0.80392857, 0.82482143, 0.8
2276786]),
'split4_test_score': array([0.81410714, 0.60080357, 0.80178571, 0.82535714, 0.8
19375 ]),
'mean_test_score': array([0.81507143, 0.59548214, 0.80135714, 0.82416071, 0.818
48214]),
'std_test_score': array([0.00310499, 0.00415638, 0.00486206, 0.00220056, 0.0034
6033]),
'rank_test_score': array([3, 5, 4, 1, 2], dtype=int32)}

```

```
In [115... predicted_dct=clf_dct.predict(X_test_pca)
```

```
In [122... print(classification_report(y_test,predicted_dct))
```

	precision	recall	f1-score	support
0	0.89	0.89	0.89	1362
1	0.95	0.95	0.95	1624
2	0.83	0.82	0.83	1390
3	0.80	0.77	0.78	1415
4	0.79	0.82	0.81	1361
5	0.76	0.74	0.75	1271
6	0.89	0.89	0.89	1391
7	0.83	0.86	0.84	1413
8	0.74	0.75	0.75	1372
9	0.79	0.78	0.78	1401
accuracy			0.83	14000
macro avg	0.83	0.83	0.83	14000
weighted avg	0.83	0.83	0.83	14000

```
In [128... confusion_matrix(y_test,predicted_dct)
```

```

Out[128... array([[1208,  2,  24,  15,  10,  35,  24,  12,  23,  9],
 [  1, 1542,  10,  8,  12,  7,  8,  13,  19,  4],
 [ 24,  14, 1140,  53,  27,  16,  35,  26,  43, 12],
 [ 23,  20,  53, 1086,  11,  81,  6,  29,  88, 18],
 [ 13,  7,  19,  4, 1115,  29,  18,  44,  27, 85],
 [ 30,  13,  15,  69,  24,  944,  40,  16,  95, 25],
 [ 23,  4,  32,  4,  29,  22, 1244,  3,  20, 10],
 [  8,  6,  19,  17,  34,  9,  4, 1209,  19, 88],
 [ 17,  19,  40,  74,  32,  77,  21,  16, 1035, 41],
 [ 14,  4,  15,  26, 112,  21,  1,  91,  29, 1088]])

```



```
In [117... with open('model_dct.pkl','wb') as f:
            pickle.dump(clf_dct,f)
```

```
In [157... lsvc=LinearSVC(dual=False)
clf_lsvc=GridSearchCV(lsvc,search_space_lin_svc,n_jobs=-1,cv=2)
clf_lsvc.fit(X_train,y_train)
```

```
Out[157... GridSearchCV(cv=2, estimator=LinearSVC(dual=False), n_jobs=-1,
              param_grid={'C': [0.001, 0.01, 0.1, 1]})
```

```
In [158... clf_lsvc.best_params_
```

```
Out[158... {'C': 0.01}
```

```
In [159... clf_lsvc.cv_results_
```

```
Out[159... {'mean_fit_time': array([ 161.88624752, 1146.01064491, 1940.61772645, 3071.01824
355]),
'std_fit_time': array([ 5.47511542, 15.340451 , 13.37943447, 27.57126641]),
'mean_score_time': array([0.07130957, 0.06468999, 0.10819721, 0.10718298]),
'std_score_time': array([0.00231433, 0.00124609, 0.00077796, 0.00109005]),
'param_C': masked_array(data=[0.001, 0.01, 0.1, 1],
                        mask=[False, False, False, False],
                        fill_value='?',
                        dtype=object),
'params': [{'C': 0.001}, {'C': 0.01}, {'C': 0.1}, {'C': 1}],
'split0_test_score': array([0.89728571, 0.90539286, 0.90496429, 0.90235714]),
'split1_test_score': array([0.89375 , 0.90242857, 0.90132143, 0.89907143]),
'mean_test_score': array([0.89551786, 0.90391071, 0.90314286, 0.90071429]),
'std_test_score': array([0.00176786, 0.00148214, 0.00182143, 0.00164286]),
'rank_test_score': array([4, 1, 2, 3], dtype=int32)}
```

```
In [160... predicted_lsvc=clf_lsvc.predict(X_test)
```

```
In [161... print(classification_report(y_test,predicted_lsvc))
```

	precision	recall	f1-score	support
0	0.94	0.97	0.96	1362
1	0.94	0.97	0.95	1624
2	0.92	0.88	0.90	1390
3	0.89	0.88	0.89	1415
4	0.91	0.94	0.92	1361
5	0.90	0.85	0.87	1271
6	0.93	0.95	0.94	1391
7	0.92	0.95	0.94	1413
8	0.88	0.86	0.87	1372
9	0.90	0.88	0.89	1401
accuracy			0.91	14000
macro avg	0.91	0.91	0.91	14000
weighted avg	0.91	0.91	0.91	14000

```
In [162... confusion_matrix(y_test,predicted_lsvc)
```

```
Out[162...] array([[1327, 1, 3, 2, 3, 5, 11, 3, 4, 3],
      [ 1, 1577, 6, 7, 2, 5, 5, 6, 14, 1],
      [ 12, 21, 1229, 22, 15, 9, 14, 20, 41, 7],
      [ 10, 13, 40, 1245, 3, 31, 8, 9, 37, 19],
      [ 3, 5, 7, 3, 1278, 4, 13, 4, 5, 39],
      [ 18, 4, 10, 53, 21, 1075, 34, 4, 39, 13],
      [ 14, 5, 13, 0, 10, 15, 1321, 4, 8, 1],
      [ 0, 6, 9, 4, 9, 2, 4, 1339, 4, 36],
      [ 13, 42, 11, 33, 12, 43, 12, 8, 1176, 22],
      [ 7, 9, 6, 26, 51, 10, 0, 52, 9, 1231]])
```

```
In [163...] with open('model_lsvc.pkl', 'wb') as f:
            pickle.dump(clf_lsvc, f)
```

```
In [166...] svc=SVC(C=0.001, max_iter=2000)
            clf_svc=GridSearchCV(svc, search_space_s, n_jobs=-1, cv=5)
            clf_svc.fit(X_train_pca, y_train)
```

/opt/anaconda3/lib/python3.8/site-packages/sklearn/svm/_base.py:255: Convergence Warning: Solver terminated early (max_iter=2000). Consider pre-processing your data with StandardScaler or MinMaxScaler.

warnings.warn('Solver terminated early (max_iter=%i).')

```
Out[166...] GridSearchCV(cv=5, estimator=SVC(C=0.001, max_iter=2000), n_jobs=-1,
                        param_grid={'kernel': ['linear', 'poly', 'rbf', 'sigmoid']})
```

```
In [167...] clf_svc.best_params_
```

```
Out[167...] {'kernel': 'linear'}
```

```
In [168...] clf_svc.cv_results_
```

```
Out[168...] {'mean_fit_time': array([ 171.95512977,  984.58973093, 1038.45634322,  751.77450
013]),
             'std_fit_time': array([ 3.13967254,  6.40468162, 19.94938339,  8.52851972]),
             'mean_score_time': array([ 27.29575744,  59.96360068, 110.68338299,  49.6229634
3]),
             'std_score_time': array([0.26372807, 0.38421839, 4.67408844, 2.68643461]),
             'param_kernel': masked_array(data=['linear', 'poly', 'rbf', 'sigmoid'],
                                           mask=[False, False, False, False],
                                           fill_value='?',
                                           dtype=object),
             'params': [{'kernel': 'linear'},
                        {'kernel': 'poly'},
                        {'kernel': 'rbf'},
                        {'kernel': 'sigmoid'}],
             'split0_test_score': array([0.93633929, 0.09928571, 0.63946429, 0.59705357]),
             'split1_test_score': array([0.93892857, 0.09991071, 0.62598214, 0.59776786]),
             'split2_test_score': array([0.94151786, 0.09946429, 0.63098214, 0.59955357]),
             'split3_test_score': array([0.93464286, 0.09955357, 0.639375, 0.60741071]),
             'split4_test_score': array([0.93857143, 0.09946429, 0.63017857, 0.61330357]),
             'mean_test_score': array([0.938, 0.09953571, 0.63319643, 0.60301786]),
             'std_test_score': array([0.00234915, 0.00020671, 0.00535756, 0.00632818]),
             'rank_test_score': array([1, 4, 2, 3], dtype=int32)}
```

```
In [169...] predicted_svc=clf_svc.predict(X_test_pca)
```

In [170...

```
print(classification_report(y_test,predicted_svc))
```

	precision	recall	f1-score	support
0	0.96	0.99	0.97	1362
1	0.96	0.98	0.97	1624
2	0.94	0.93	0.93	1390
3	0.93	0.91	0.92	1415
4	0.94	0.96	0.95	1361
5	0.92	0.90	0.91	1271
6	0.96	0.96	0.96	1391
7	0.95	0.95	0.95	1413
8	0.93	0.92	0.92	1372
9	0.94	0.92	0.93	1401
accuracy			0.94	14000
macro avg	0.94	0.94	0.94	14000
weighted avg	0.94	0.94	0.94	14000

In [171...

```
confusion_matrix(y_test,predicted_svc)
```

```
Out[171...] array([[1342,  0,  2,  1,  3,  2,  7,  0,  5,  0],
       [  1, 1586,  8,  5,  3,  1,  4,  6,  8,  2],
       [ 15,  10, 1298, 12, 12,  3, 11, 11, 16,  2],
       [  3,  8,  31, 1285,  2, 38,  4,  9, 29,  6],
       [  3,  3, 10,  0, 1313,  1,  5,  2,  2, 22],
       [ 11,  7,  5, 39,  7, 1149, 20,  1, 26,  6],
       [ 13,  3,  7,  1,  7,  15, 1341,  1,  3,  0],
       [  0,  5, 10,  5, 10,  0,  0, 1346,  1, 36],
       [  8, 25,  9, 19,  8, 29,  6,  6, 1257,  5],
       [  3,  6,  7, 19, 35,  5,  0, 33,  8, 1285]])
```

In [172...

```
with open('model_svc.pkl','wb') as f:
    pickle.dump(clf_svc,f)
```

In []: